

Lab 2 — Sprint Plan, Estimation & Reflection

Restaurant Table Booking & Pre-Ordering App | Problem Statement #36

PES University — Dept. of CSE • Name: _____ SRN: _____ Date: _____

1. Sprint 1

Sprint goal: A diner can see the live floor plan and hold a specific table for a chosen dining window.

Duration: 1 week • Committed: 18 story points • Completed: 13 • Carried over: 5 (Story 1.2)

ID	User story	Epic	Priority	SP	End-of-sprint status
1.1	View live floor plan	Epic 1	Highest	8	Done
1.2	Filter tables by slot and party size	Epic 1	High	5	Not completed → carried over
1.3	Select and hold a table	Epic 1	Highest	5	Done

2. Sprint 2

Sprint goal: Double booking is impossible and a diner can release a table they no longer need.

Duration: 1 week • Committed: 11 story points • Completed: 11 • Story 1.2 stayed in the backlog rather than being re-committed

ID	User story	Epic	Priority	SP	End-of-sprint status
1.2	Filter tables by slot and party size (carried over)	Epic 1	High	5	Not completed → carried over
1.4	Reject overlapping reservations	Epic 1	Highest	8	Done
1.5	Cancel a reservation	Epic 1	Medium	3	Done

3. Velocity

Sprint	Duration	Committed	Completed	Achieved	Outcome
Sprint 1	1 week	18	13	72 %	Story 1.2 not completed, returned to the backlog
Sprint 2	1 week	11	11	100 %	All committed work accepted; Story 1.2 not pulled back in
Average	1 week	14.5	12	—	Working velocity ≈ 12 pts per week; 10 pts remain

4. Planning Poker record

Each story was discussed briefly, three estimators voted at the same time, and outliers explained their reasoning before a consensus number was agreed. The five stories below produced the widest spread.

Story	Votes	Agreed	What the discussion surfaced
1.1 View live floor plan	8 / 5 / 13	8	Rendering a live, colour-coded plan is more than a static image — the 13 vote flagged real-time refresh, which the team agreed belongs to a separate sync story.
1.2 Filter tables by slot and party size	3 / 5 / 5	5	Everyone priced the filter itself. Nobody priced wiring it to the live availability service, which is why it was the story that slipped.
1.4 Reject overlapping reservations	5 / 8 / 8	8	Concurrency control is the highest-risk item in the epic; the team took the higher estimate because of uncertainty, not volume of code.
1.5 Cancel a reservation	3 / 3 / 5	3	State change plus a cutoff rule over data that already exists — quick agreement.

5. Burndown analysis

See [Lab2_Burndown_Charts.pdf](#) for both sprints, and the Jira exports in screenshots/. The grey guideline is where the team should be assuming linear progress; the red line is the work actually remaining. Sprint 1 steps down twice and then flattens at 5 points, because Story 1.2 was never accepted. Sprint 2 steps down cleanly to zero. The Jira exports in Appendix B differ in two ways worth explaining: the Sprint 1 line spikes to 29 because stories were moved in and out after the sprint had started, and it finishes at zero rather than 5 because Jira removes the incomplete story from the sprint at the moment the sprint is completed. Because both sprints were run in one session, the Jira charts compress every step onto the first day; the sequence of values is still the real one.

6. Reflection

1. Did your estimations reflect the actual effort?

Partly. The two stories sized for risk rather than volume — 1.1 View live floor plan at 8 points, and 1.4 Reject overlapping reservations, also 8 — were both accepted without trouble, which suggests the team was right to price uncertainty rather than lines of code. The miss was Story 1.2 Filter tables by slot and party size. Planning Poker produced 3, 5, 5 with almost no discussion, because every estimator priced the filter itself and nobody priced wiring it to the live availability service. The quick agreement was the warning sign, not the reassurance it felt like at the time, and it is the one story that did not finish.

2. Was your backlog well-prioritized?

Within the epic, yes. Both Highest stories that make a booking trustworthy — 1.3 Select and hold a table and 1.4 Reject overlapping reservations — were scheduled ahead of the Medium cancellation and no-show work, so the risk of two parties being sent to the same table closed before convenience features were touched. The mistake shows up between the sprints: Story 1.2 was left in the backlog rather than re-committed at the top of Sprint 2. Unfinished work should normally be pulled forward ahead of new stories, otherwise a half-built feature ages while newer work overtakes it. Deferring 1.6 Auto-release a no-show table was the right call — it recovers revenue but nothing depends on it.

3. How did your simulated sprint align with your plan?


Less cleanly than the plan reads, and the Jira burndown shows exactly why. Sprint 1 was started before its contents were correct, so stories were added and removed after the sprint was already running; the remaining-values line spikes to 29 points before settling at the intended 18. The guideline is anchored to the 11 points the sprint held at the moment it started, which is why it does not begin at 18 — a direct illustration of why sprint scope is fixed at planning. From that point the sprint ran as planned: 18 points committed, 13 accepted (Stories 1.1 and 1.3), and Story 1.2 left in progress. Sprint 2 committed 11 points and accepted all of them. Both sprints were simulated in a single session, so the charts compress every step onto the first day.

4. What insights did the burndown chart give about your team's capacity?


The team accepted 13 points and then 11, so working capacity is around 12 points per week against an average commitment of 14.5 — a consistent over-commitment of roughly 20 %. With 10 points still outstanding (Story 1.2 at 5 and Story 1.6 at 5), the epic needs one more short sprint to close. The second lesson is about story size: in Sprint 1 the line moves twice and then flattens, so from the midpoint onwards the chart showed no progress even though work was happening. An 8-point story hides its state until it is accepted. The third lesson is procedural — changing sprint contents after the sprint has started makes the guideline meaningless for the rest of the sprint, because it is calculated from the commitment at start.

7. Appendix B — Jira evidence

B1. Sprint 1 active board — work spread across To Do, In Progress and Done



B2. Sprint 1 burndown chart (Story Points)



B3. Sprint 2 burndown chart (Story Points)

